

2026

TECHNICAL RESEARCH REPORT

Global Landscape of LLM Evaluation Aggregators, Cost- Efficiency Trackers, and Live Benchmarking Suites

High-Density Wide-Scale Research Sweep: Agility, Reasoning, Efficiency, Agentic
Benchmarks, and REST API Integration Vectors

A comprehensive mapping of 40+ platforms across six structural dimensions: model
inclusion agility, contamination-resistant reasoning evaluation, execution efficiency
profiling, autonomous agentic tool-calling benchmarks, REST API integration vectors, and
schema-level authentication specifications.

CLASSIFICATION

Principal AI Platform Architect — System Performance Engineering

DATE

May 2026

VERSION

1.0 — Final

Table of Contents

1. Executive Summary	3
2. High-Speed Agility Trackers	3
2.1 Speed Classification	3
2.2 Verification Before Formal Validation	4
2.3 The Speed vs. Rigor Trade-off	4
3. Contamination-Resistant Reasoning Evaluation	5
3.1 Three Paradigms of Contamination Resistance	5
3.2 Contamination Resistance Taxonomy	6
3.3 Recommended "Fluid IQ" Evaluation Stack	6
3.4 Benchmarks to Deprecate	7
4. Execution Efficiency and Resource Profiling	7
4.1 Efficiency Metric Coverage	7
4.2 Critical Gaps in Efficiency Benchmarking	8
4.3 Provider-Level vs. Self-Hosted Benchmarks	8
5. Autonomous Agentic and Tool-Calling Benchmarks	9
5.1 Capability Coverage Matrix	9
5.2 MCP-Specific Benchmarks	9
5.3 Software Engineering Agent Benchmarks	10
6. REST API Integration and Data Retrieval	10
6.1 Access Tier Classification	11

6.2 Key API Endpoint Details	12
6.2.1 OpenRouter API	12
6.2.2 Artificial Analysis API	12
6.2.3 ZeroEval/LLM Stats API	12
6.2.4 Hugging Face Hub API (Generic Dataset Access)	12
7. Schema Specifications and Authentication Protocols	13
7.1 API Schema Reference	13
7.1.1 OpenRouter Schema	13
7.1.2 ZeroEval/LLM Stats Schema	13
7.1.3 Artificial Analysis Schema	14
7.2 Authentication Protocol Summary	14
8. Technical Appendix: Complete Platform Directory	15
8.1 Agility Trackers and General Aggregators	15
8.2 Contamination-Resistant Reasoning Benchmarks	15
8.3 Agentic and Tool-Calling Benchmarks	16
8.4 Efficiency and Cost Profiling Platforms	17
8.5 Recommended Data Aggregation Architecture	17

1. Executive Summary

This report presents a comprehensive mapping of the entire global landscape of real-time LLM evaluation aggregators, automated cost-efficiency trackers, and live benchmarking suites. The research decomposes the ecosystem into six structural dimensions: high-speed model inclusion agility, contamination-resistant reasoning evaluation, execution efficiency profiling, autonomous agentic tool-calling benchmarks, REST API integration vectors, and schema-level authentication specifications. The findings are drawn from primary sources including GitHub code repositories, Hugging Face Spaces, official API documentation, project whitepapers, and developer community discussions on social platforms.

Across all six dimensions, we profile over 40 distinct platforms and benchmarks. A central finding is that no single platform provides both the fastest model inclusion turnaround and comprehensive programmatic API access: these attributes are inversely correlated. The most authoritative quality ranking (LMSYS Chatbot Arena) has no public REST API, while the most API-friendly platforms (Artificial Analysis, OpenRouter) provide excellent programmatic access but aggregate rather than independently evaluate models. In the reasoning domain, the field has pivoted from static public benchmarks (MATH, BBH) to three contamination-resistance paradigms: live/dynamic refresh, private holdout evaluation, and novel-question design. For agentic evaluation, only four benchmarks test MCP (Model Context Protocol) proficiency, representing a critical gap given MCP's rapid industry adoption. In efficiency profiling, context-window degradation latency and cost-per-complex-task remain the least-covered metrics globally, with no single platform comprehensively tracking all five core efficiency parameters.

The report is organized as follows: Sections 2 through 7 address each structural question in depth, with comparative summary tables. Section 8 provides the technical appendix detailing REST API schemas, authentication protocols, and alternative programmatic fetch paths for every discovered system.

2. High-Speed Agility Trackers

This section maps platforms by their model inclusion turnaround time, defined as the elapsed time between a model's public announcement (on X/Twitter, Hugging Face, or provider blogs) and its appearance on the evaluation platform with at least preliminary scores. We analyze operational infrastructure, verification mechanisms, and API access for each platform.

2.1 Speed Classification

Platforms cluster into three speed tiers based on inclusion turnaround. The fastest tier achieves same-day to 48-hour inclusion through crowdsourced voting, automated API probing, or real-time production traffic. The moderate tier takes one to seven days, typically requiring independent benchmark execution or waiting for aggregated results from primary sources. The slowest tier takes weeks to months, due to academic rigor, compute constraints, or manual expert evaluation processes.

Platform	Inclusion Speed	Verification	Infrastructure	API Access
LMSYS Chatbot Arena	Same day - 48h	Crowdsourced ELO	Crowdsourced	No public API
Artificial Analysis	Same day - 48h	Independent auto benchmarks	Automated/hybrid	Free REST API
LLM Stats	Hours - 1 day	Aggregation + internal	Automated, hourly	Not documented
OpenRouter Rankings	Same day	Usage + benchmark agg.	Production traffic	Full REST API
Aider Leaderboard	1 - 7 days	Automated, deterministic	Automated	No (OSS tool)
LiveBench	1 - 3 days	Automated, contamination-free	Automated	GitHub only
Vellum Leaderboard	1 - 7 days	Aggregation only	Automated aggregation	No
WhatLLM	Weekly	Aggregation (AA Index)	Automated aggregation	No
HF Open LLM Leaderboard	Days - weeks	Automated (lighteval)	Queue-based	HF datasets API
Scale SEAL	Weeks - months	Expert human + private data	Human expert	No
Stanford HELM	Weeks - months	Academic, reproducible	Academic/manual	Python framework

Table 1: Platform Agility Comparison Matrix

2.2 Verification Before Formal Validation

Most fast-inclusion platforms use provisional scoring mechanisms to balance speed with reliability. LMSYS Chatbot Arena allows models to appear within hours, but requires approximately 1,000+ crowdsource votes for a statistically stable Elo rating. New models start with provisional ratings based on Bayesian priors, and confidence intervals are computed and displayed alongside scores. Artificial Analysis runs independent benchmark suites with 10+ repeats to compute 95% confidence intervals, publishing results only once statistical significance is achieved. The Hugging Face Open LLM Leaderboard takes a different approach: no scores appear until the full automated evaluation completes on their GPU cluster, making it slower but fully verified at publication time. Notably, no platform currently implements a formal "preview/preliminary" flag system that clearly distinguishes unverified from verified scores in a machine-readable way.

2.3 The Speed vs. Rigor Trade-off

There is a clear inverse relationship between inclusion speed and evaluation rigor. The fastest platforms (LMSYS, Artificial Analysis, LLM Stats) achieve speed through crowdsourcing, automation, or aggregation, but sacrifice depth of evaluation. The most rigorous platforms (HELM, Scale SEAL) require weeks to months but provide the most thorough and trustworthy evaluations. The sweet spot is occupied by Aider and LiveBench, which provide automated, reproducible evaluations within days. For building an agility monitoring system, the recommended primary signal sources are: OpenRouter API (real-time model listings), Artificial Analysis API (quality + speed + price data), and the community-maintained LMSYS Arena scraper at github.com/oolong-tea-2026/arena-leaderboard-archive.

3. Contamination-Resistant Reasoning Evaluation

The benchmarking landscape has undergone a fundamental shift. Static benchmarks like MMLU and the original MATH dataset are widely considered compromised by data contamination, with frontier models achieving near-ceiling scores that reflect memorization rather than genuine reasoning. This section documents 14 active benchmarks across three contamination-resistance paradigms: Live/Dynamic Refresh, Private Holdout Evaluation, and Novel-Question Design.

3.1 Three Paradigms of Contamination Resistance

The Live/Dynamic Refresh paradigm continuously generates new evaluation questions from fresh sources, ensuring any contamination window is narrow. LiveBench refreshes questions monthly from newly-published arXiv papers, Codeforces problems, and competition mathematics. LiveCodeBench harvests competitive programming problems with precise release-date tagging, enabling time-stratified evaluation that directly measures contamination by comparing pre-cutoff and post-cutoff performance. MathArena evaluates models on math competition problems immediately after release, before they can enter training data, and has provided direct evidence of contamination in AIME 2024 results. SWE-rebench maintains an automated pipeline that continuously mines fresh GitHub issues, producing over 21,000 interactive software engineering tasks.

The Private Holdout paradigm never releases evaluation data publicly, making contamination impossible by construction. ARC-AGI-2 employs a three-tier evaluation structure: public eval (120 tasks), semi-private eval (120 tasks available via API for remote testing only), and private eval (never released, used for official Kaggle competition scoring). Scale AI SEAL uses fully private datasets curated by verified domain experts, with human grading across coding, math, instruction following, and conversational evaluation. SWE-bench Pro applies license restrictions to evaluation code as a legal contamination deterrent, combined with private task curation by Scale AI. These approaches offer the strongest contamination resistance but sacrifice reproducibility and transparency, requiring the community to trust the platform's methodology.

The Novel-Question Design paradigm creates questions that require genuine reasoning unlikely to appear in training data. ARC-AGI tasks use grid-based pattern induction that is inherently resistant to

memorization because each task requires novel pattern recognition. GPQA (Google-Proof Q&A;) presents 448 graduate-level multiple-choice questions designed so that even with internet access, non-experts cannot answer them. Humanity's Last Exam (HLE) features 2,500 questions authored by 1,000+ domain experts at the frontier of human knowledge. However, the resistance of all novel-design benchmarks degrades once questions become public, as they can then enter future training data. HLE acknowledges this by positioning itself as a one-time snapshot requiring successor exams.

3.2 Contamination Resistance Taxonomy

Benchmark	Resistance Mechanism	Strength	Refresh Rate	Reasoning Focus
LiveBench	Monthly question refresh	4/5	Monthly	General (6 categories)
HLE	Novel expert questions	3/5 (5/5 at launch)	One-time	Broad academic
ARC-AGI-2	Private eval + novel design	5/5	Versioned (years)	Fluid intelligence
GPQA Diamond	Expert-authored, Google-proof	3/5 (public now)	Static	Scientific reasoning
SWE-rebench	Continuous automated mining	5/5	Continuous	Software engineering
LiveCodeBench	Date-tagged contest problems	5/5	Continuous	Code generation
SEAL Leaderboards	Fully private datasets	5/5	Periodic	Multi-domain
MathArena	Live competition tracking	5/5	Per-competition	Mathematical reasoning
BBH / BBEH	None (static public)	1/5	Static	Multi-step reasoning
MATH / Minerva	None (contaminated)	1/5	Static	Mathematical reasoning

Table 2: Contamination Resistance Taxonomy

3.3 Recommended "Fluid IQ" Evaluation Stack

For isolating authentic reasoning from memorization, the recommended evaluation stack consists of five benchmarks covering different contamination-resistance paradigms and reasoning domains. First, ARC-AGI-2 (private eval) serves as the gold standard for fluid intelligence measurement, as it was explicitly designed to measure novel problem-solving versus memorized knowledge. Second, MathArena provides the most rigorous live math evaluation, offering direct contamination evidence through its timestamped problem tracking. Third, LiveCodeBench delivers the most rigorous contamination measurement for code reasoning through its date-tagged competitive programming problems. Fourth, LiveBench provides broad contamination-limited evaluation across six categories

with monthly question refresh. Fifth, Scale SEAL private benchmarks offer the most commercially trusted private evaluation, particularly for enterprise decision-making where reproducibility is less critical than trustworthiness.

3.4 Benchmarks to Deprecate

Several widely-cited benchmarks should no longer be used as primary evaluation tools for frontier models. MATH and MATH-500 are saturated with scores exceeding 95% for frontier models and are widely contaminated. SWE-bench Verified has been officially deprecated by OpenAI, which identified at least 16.4% flawed test cases and evidence of training data contamination. BIG-Bench Hard (BBH) is saturated for frontier models. IFEval achieves near-ceiling performance for most models and measures instruction compliance rather than reasoning depth. AIME 2024 has been demonstrated as significantly contaminated by the MathArena research team. The trajectory of SWE-bench illustrates the full contamination lifecycle: gold standard in 2023, human-filtered subset in 2024, official deprecation in 2025, and ecosystem migration to SWE-rebench and SWE-bench Pro.

4. Execution Efficiency and Resource Profiling

This section maps all active evaluation suites that explicitly track model execution efficiency parameters: Tokens Per Second (TPS), Time to First Token (TTFT), context-window degradation latency, blended API cost-per-million tokens, and cost-per-complex-task calculations. We identify 15+ platforms and highlight critical gaps in the current landscape.

4.1 Efficiency Metric Coverage

Platform	TPS	TTFT	Context Degr.	Cost/1M Tokens	Cost-per-Task	API Access
Artificial Analysis	Yes	Yes	No	Yes	No	REST API
HF LLM-Perf Leaderboard	Yes	Yes	No	No	No	HF datasets
OpenRouter	No	No	No	Yes	Partial	REST API
Anyscale/Ray Serve	Yes	Yes	Partial	No	No	No
vLLM Benchmark Suite	Yes	Yes	Partial	No	No	No
TensorRT-LLM/GenAI-Perf	Yes	Yes	Partial	Partial	No	No
Stanford HELM	Partial	No	No	No	No	No
AI21 Maestro	No	Yes	No	No	Yes	No

SemiAnalysis InferenceMAX	Yes	Yes	Partial	Yes	Partial	No
MLPerf Inference	Yes	Yes	No	No	No	No
Epoch AI	No	No	No	Yes	Yes	No
Modal LLM Almanac	Yes	Yes	No	Yes	Partial	No
BentoML llm-optimizer	Yes	Yes	No	Partial	No	No
AIMultiple	Yes	Yes	Yes	Yes	No	No

Table 3: Efficiency Metric Coverage Matrix

4.2 Critical Gaps in Efficiency Benchmarking

No single platform covers all five target metric categories comprehensively. The most comprehensive platforms (Artificial Analysis, SemiAnalysis InferenceMAX) each cover only three to four of the five categories. Two metrics stand out as critically under-served. Context-window degradation latency, which measures how TTFT and throughput degrade as input context length increases toward the model's maximum, is tracked by only AIMultiple (via their 32-message conversation test) and partially by Anyscale, vLLM, and NVIDIA through configurable input sequence length sweeps. No platform publishes systematic latency-versus-context-length curves for major API providers. Cost-per-complex-task, which captures the real cost of multi-turn, reasoning, or agentic workflows beyond simple per-token pricing, is addressed by only AI21 Maestro (for agentic tasks), Epoch AI (price-for-fixed-performance tracking), and Digital Applied (blog post estimates). Per-million-token pricing does not capture the real cost of complex workflows that may require multiple API calls, retries, and extended reasoning chains.

4.3 Provider-Level vs. Self-Hosted Benchmarks

The efficiency landscape is further divided between provider-level benchmarks (measuring API endpoints from commercial providers like OpenAI, Anthropic, and Google) and self-hosted benchmarks (measuring open-weight models on specific hardware configurations). Provider-level benchmarks include Artificial Analysis, OpenRouter, LMSYS, HELM, and Epoch AI. Self-hosted benchmarks include vLLM, TensorRT-LLM, Anyscale, MLPerf, InferenceMAX, HF LLM-Perf, BentoML, and Modal. These two dimensions are rarely combined on the same platform. Only Artificial Analysis attempts to bridge the gap by comparing provider endpoints, but it cannot measure self-hosted performance. This means that a direct comparison of "GPT-5 via API" versus "Llama-4-70B on H100 via vLLM" on the same efficiency axes remains impossible without custom integration work. SemiAnalysis InferenceMAX is the fastest-updating benchmark in the landscape, running nightly automated benchmarks on real hardware across NVIDIA H100, B200, and AMD MI300X GPUs with vLLM, TensorRT-LLM, and SGLang serving engines. It uniquely tracks energy efficiency via tokens-per-megawatt metrics.

5. Autonomous Agentic and Tool-Calling Benchmarks

This section maps 22 benchmarks that explicitly measure long-horizon agentic behavior, deep multi-tool execution pipelines, and native Model Context Protocol (MCP) tool usage proficiency. The landscape divides into five functional categories: Software Engineering Agents, Web/Desktop Environment Agents, General Tool-Use and Function Calling, Multi-Environment/Multi-Turn Evaluation, and MCP-Specific Benchmarks. A critical finding is that only four benchmarks test MCP proficiency, representing a significant gap given MCP's rapid industry adoption by Anthropic, OpenAI, and other major providers.

5.1 Capability Coverage Matrix

Benchmark	Tool Use	Multi-Step	Error Recovery	Long-Horizon	MCP	Multi-Turn
SWE-bench	Yes	Yes	Yes (test-driven)	Yes	No	Yes
Terminal-Bench	Yes	Yes	Yes	Yes	No	Yes
MCP-Atlas	Yes	Yes	No	No (single-turn)	Yes	No
WebArena	Yes	Yes	Yes	Yes	No	Yes
OSWorld	Yes	Yes	Yes	Yes	No	Yes
GAIA	Yes	Yes	Yes	Yes	No	Yes
AgentBench	Yes	Yes	Yes	Yes	No	Yes
BFCL v4	Yes	Partial	No	No	No	Partial
ToolBench	Yes	Yes	Partial	Yes	No	Yes
tau-bench/tau2-bench	Yes	Yes	Yes	Yes	No	Yes
MCP-Bench	Yes	Yes	Partial	Yes	Yes	Yes
MCP-AgentBench	Yes	Yes	Partial	Partial	Yes	Partial
MCPVerse	Yes	Yes	Partial	Yes	Yes	Yes
LongCLI-Bench	Yes	Yes	Yes	Yes	No	Yes

Table 4: Agentic Capability Coverage Matrix

5.2 MCP-Specific Benchmarks

The Model Context Protocol (MCP) has rapidly emerged as the industry standard for tool-server integration, adopted by Anthropic, OpenAI, and other major providers. Yet only four benchmarks

explicitly test MCP proficiency. MCP-Atlas is the first large-scale benchmark designed for MCP evaluation, featuring 36 real MCP servers with 220 tools and 1,000 realistic single-turn tasks. It tests tool discovery, selection, and orchestration over live MCP servers, using AST-based matching for function call correctness. MCP-Bench, developed by Accenture, evaluates LLM agents through the MCP protocol with multi-step task execution and cross-tool coordination. MCP-AgentBench, published at AAAI 2025, is the first benchmark to rigorously evaluate language agents operating within the MCP ecosystem, focusing on MCP as the integration layer between agents and tool servers. MCPVerse is the largest MCP benchmark by tool count, featuring 550+ real-world executable tools creating an action space exceeding 140,000 tokens.

A critical gap exists: no benchmark comprehensively tests long-horizon planning (10+ sequential tool calls), error recovery from MCP server failures, and MCP proficiency simultaneously. MCP-Atlas is single-turn only. MCPVerse tests multi-step execution but lacks step-level error recovery evaluation. LongCLI-Bench is the only benchmark explicitly designed for long-horizon agentic programming, finding that the majority of tasks stall at less than 30% completion, but it does not use MCP protocol. The recommended next action is to build an MCP-native long-horizon benchmark that tests 10-50 sequential MCP tool calls, error recovery from server failures, cross-server orchestration, and step-level progress tracking.

5.3 Software Engineering Agent Benchmarks

The software engineering agent benchmark ecosystem is the most mature and rapidly evolving segment of the agentic evaluation landscape. SWE-bench, the de facto standard since 2023, evaluates agents on real GitHub issues using Docker-based test harnesses with execution verification. However, as documented in Section 3, SWE-bench Verified has been deprecated by OpenAI due to contamination and flawed test cases. The ecosystem is migrating to three successors: SWE-rebench (continuous automated GitHub mining, 21,000+ tasks), SWE-bench Pro (Scale AI's longer-horizon tasks with license restrictions), and Multi-SWE-bench (multilingual coverage of Java, TypeScript, Go, Rust, C, and C++ with 2,132 instances). Terminal-Bench 2.0 provides 89 terminal-based tasks requiring end-to-end execution in real Docker environments, with milestone-level tracking beyond binary pass/fail. FeatureBench, published at ICLR 2026, is the first benchmark for feature-level coding (versus bug-fixing in SWE-bench), requiring agents to implement complete features that pass automated test suites.

6. REST API Integration and Data Retrieval

This section catalogs every discovered programmatic access vector for LLM evaluation data across all platforms profiled in Sections 2 through 5. Of the 17+ platforms investigated, only four have fully documented, production-grade REST APIs for leaderboard and benchmark data retrieval. The remaining platforms expose data through Hugging Face datasets (Parquet format), GitHub repositories, or have no programmatic access beyond web scraping.

6.1 Access Tier Classification

Tier	Platform	Access Method	Auth Required
Tier 1: Full REST API	OpenRouter	/api/v1/models (OpenAI-compatible)	API key (chat); none (models list)
Tier 1: Full REST API	Artificial Analysis	/api/* (x-api-key header)	API key (1K req/day free)
Tier 1: Full REST API	ZeroEval/LLM Stats	api.zeroeval.com/v1/*	API key
Tier 2: HF Hub API	Open LLM Leaderboard	open-llm-leaderboard/results dataset	HF token
Tier 2: HF Hub API	SWE-bench	princeton-nlp/SWE-bench_Verified	HF token
Tier 2: HF Hub API	GPQA	Idavidrein/gpqa (gated)	HF token + agreement
Tier 2: HF Hub API	GAIA	gaia-benchmark/GAIA (gated)	HF token + agreement
Tier 2: HF Hub API	BFCL	gorilla-llm/Berkeley-Function-Calling-Leaderboard	HF token
Tier 2: HF Hub API	LMSYS Arena	lmsys/chatbot_arena_conversations	HF token
Tier 2: HF Hub API	ToolBench	tuandunghcmut/toolbench-v1	None
Tier 3: GitHub only	ARC-AGI	arcprize/ARC-AGI-Community-Leaderboard	None
Tier 3: GitHub only	AgentBench	THUDM/AgentBench	None
Tier 3: GitHub only	WebArena	web-arena-x/webarena	None
Tier 3: GitHub only	OSWorld	xlang-ai/osworld	None
Tier 3: GitHub only	Terminal-Bench	harbor-framework/terminal-bench	None
Tier 3: GitHub only	HELM	stanford-crfm/helm	None
Tier 4: No access	Scale AI SEAL	labs.scale.com (web only)	None
Tier 4: No access	LiveBench (website)	livebench.ai (web only)	None

Table 5: Programmatic Access Tier Classification

6.2 Key API Endpoint Details

6.2.1 OpenRouter API

OpenRouter provides the most accessible API in the landscape. The models list endpoint requires no authentication and returns 358+ models with real-time pricing, context window specifications, architecture details, and supported parameters. The base URL is <https://openrouter.ai/api/v1> with the primary endpoint `GET /api/v1/models`. The response schema includes per-model fields: `id`, `name`, `context_length`, `architecture` (modality, input/output modalities, tokenizer), `pricing` (prompt, completion, input_cache_write costs per token), `top_provider` (`context_length`, `max_completion_tokens`, `moderation status`), and `supported_parameters` (include `reasoning`, `logprobs`, `tools`, `structured_outputs`, etc.). Query parameter `?supported_parameters=json_schema` filters models by capability. The chat completions endpoint `POST /api/v1/chat/completions` is OpenAI-compatible and requires an API key for authentication.

6.2.2 Artificial Analysis API

Artificial Analysis offers a comprehensive REST API at <https://artificialanalysis.ai/api/> with API key authentication via the `x-api-key` header. The free tier provides 1,000 requests per day. Key endpoints include `GET /api/models` (model metadata and intelligence scores), `GET /api/pricing` (pricing data across providers), `GET /api/speed` (throughput and latency benchmarks), `GET /api/elo` (Elo ratings including image editing models), and various per-benchmark evaluation endpoints (e.g., `/api/evaluations/livecodebench`). The response schema includes model-level data with quality metrics, speed benchmarks (output tokens/sec, TTFT), and pricing per million tokens across multiple providers. An MCP skill wrapper is also available at smithery.ai/skills/alexfazio/artificial-analysis-compare for direct integration with AI agent workflows.

6.2.3 ZeroEval/LLM Stats API

The ZeroEval API at <https://api.zeroeval.com/v1/> provides model listings with pricing, capabilities, and provider details. The `GET /v1/models` endpoint returns an array of model objects with fields including `id`, `display_name`, `provider_name`, `organization_name`, `context_length`, `input_price`, `output_price` (in cents per million tokens), `quantization_type`, `input/output modalities`, `routing_providers`, and `tier classification`. The `GET /v1/models/{id}` endpoint provides detailed model information including per-provider specifications with feature flags (`function_calling`, `structured_output`, `batch_inference`, `finetuning`). An MCP server is available for programmatic access. The benchmark scores endpoint was not found at `/v1/scores` or `/v1/leaderboard` during testing, suggesting it may require additional authentication or use different endpoint paths.

6.2.4 Hugging Face Hub API (Generic Dataset Access)

The Hugging Face Hub API at <https://huggingface.co/api/> serves as the universal access vector for any benchmark published as a HF dataset. Key endpoints include `GET /api/datasets/{dataset_id}` (dataset metadata), `GET /api/datasets/{dataset_id}/parquet` (Parquet export), and direct file download via dataset URLs. Authentication uses HF tokens, which are required for gated or private datasets (GPQA, GAIA).

The Python datasets library (pip install datasets) provides programmatic access with `load_dataset("DATASET_ID", token="YOUR_HF_TOKEN")`. This is the recommended approach for batch-refresh data pipelines, as it handles caching, streaming, and authentication automatically.

7. Schema Specifications and Authentication Protocols

This section provides exact base URLs, query parameters, authentication configurations, and structural JSON response schemas for all discovered programmatic access points. The information is organized as a technical reference for building an automated parsing engine that can map payloads from multiple sources into a unified data model.

7.1 API Schema Reference

7.1.1 OpenRouter Schema

Field	Type	Description
data[].id	string	Model identifier (e.g., "qwen/qwen3.7-max")
data[].name	string	Display name
data[].context_length	integer	Maximum context window in tokens
data[].architecture.modality	string	Input-output modality (e.g., "text->text")
data[].pricing.prompt	string	Cost per prompt token (USD)
data[].pricing.completion	string	Cost per completion token (USD)
data[].pricing.input_cache_write	string	Cost per cached input token (USD)
data[].top_provider.context_length	integer	Provider-specific context limit
data[].top_provider.max_completion_tokens	integer	Maximum output tokens
data[].supported_parameters	string[]	Capabilities: tools, structured_outputs, etc.

Table 6: OpenRouter /api/v1/models Response Schema

7.1.2 ZeroEval/LLM Stats Schema

Field	Type	Description
[].id	string	Model identifier (e.g., "claude-opus-4-6")
[].display_name	string	Human-readable model name
[].organization_name	string	Model provider organization

[].context_length	integer	Maximum context window in tokens
[].input_price	float	Input price in cents per million tokens
[].output_price	float	Output price in cents per million tokens
[].input_modalities	string[]	Supported input types (text, image)
[].routing_providers	string[]	Available routing providers
[].tier	string	Model tier classification

Table 7: ZeroEval /v1/models Response Schema

7.1.3 Artificial Analysis Schema

Field	Type	Description
model_name	string	Model identifier
intelligence_index	float	Composite quality score
output_speed	float	Tokens per second
ttft	float	Time to first token in seconds
pricing.input_per_1m	float	Input price per 1M tokens (USD)
pricing.output_per_1m	float	Output price per 1M tokens (USD)
context_window	integer	Maximum context length
elo_rating	float	Arena Elo rating

Table 8: Artificial Analysis /api/models Response Schema

7.2 Authentication Protocol Summary

Platform	Auth Method	Credential Source	Rate Limits	Free Tier
OpenRouter	Bearer token / API key	openrouter.ai signup	Varies by plan	Yes (limited)
Artificial Analysis	x-api-key header	artificialanalysis.ai signup	1,000 req/day	Yes
ZeroEval/LLM Stats	API key	zeroeval.com or llm-stats.com	Not documented	Yes
Hugging Face Hub	HF token	huggingface.co settings	Polite rate limit	Yes
GPQA (gated)	HF token + agreement	HF dataset gate approval	Standard HF limits	Yes
GAIA (gated)	HF token + agreement	HF dataset gate approval	Standard HF limits	Yes

LMSYS Arena	None (scraping)	N/A	N/A	N/A
Scale SEAL	None (web only)	N/A	N/A	N/A

Table 9: Authentication Protocol Summary

8. Technical Appendix: Complete Platform Directory

This appendix provides a comprehensive directory of every platform, benchmark, and aggregator discovered during this research, organized by category with full URLs, access methods, and source citations.

8.1 Agility Trackers and General Aggregators

Platform	URL	Primary Access	Key Source
LMSYS Chatbot Arena	arena.ai	HF datasets + community scraper	lmsys.org/blog/2024-03-01-policy
Artificial Analysis	artificialanalysis.ai	REST API (x-api-key)	artificialanalysis.ai/api-reference
LLM Stats	llm-stats.com	api.zerogeval.com/v1/*	llm-stats.com/developer
OpenRouter	openrouter.ai	REST API (OpenAI-compatible)	openrouter.ai/docs/api/reference/overview
Aider Leaderboard	aider.chat/docs/leaderboards	GitHub (OSS tool)	github.com/Aider-AI/aider
LiveBench	livebench.ai	GitHub repo	arxiv.org/abs/2406.19314
Vellum Leaderboard	vellum.ai/llm-leaderboard	Web only	vellum.ai/llm-leaderboard
WhatLLM	whatllm.org	Web only	whatllm.org/methodology
HF Open LLM Leaderboard	huggingface.co/spaces/open-llm-leaderboard	HF datasets API	huggingface.co/docs/hub/leaderboard-data-guide
Scale SEAL	labs.scale.com	Web only (no API)	scale.com/blog/leaderboard
Stanford HELM	crfm.stanford.edu/helm	Python framework + GitHub	arxiv.org/abs/2211.09110

Table 10: Agility Trackers and General Aggregators Directory

8.2 Contamination-Resistant Reasoning Benchmarks

Benchmark	URL	Resistance	Access Method
LiveBench	livebench.ai	Monthly refresh	GitHub: LiveBench/LiveBench

HLE	lastexam.ai	Novel expert questions	HF: lastexamai/hle
ARC-AGI-2	arcprize.org	Private eval + novel design	Kaggle + GitHub: fchollet/ARC-AGI
GPQA Diamond	github.com/idaudrein/gpqa	Expert-authored, Google-proof	HF: Idavidrein/gpqa (gated)
LiveCodeBench	livecodebench.github.io	Date-tagged contest problems	HF: livecodebench + GitHub
SWE-rebench	swe-rebench.com	Continuous automated mining	HF: SWE-rebench + arXiv:2505.20411
SWE-bench Pro	labs.scale.com	Private + license restrictions	GitHub: scaleapi/SWE-bench_Pro-os
SEAL Leaderboards	labs.scale.com	Fully private datasets	Web only (no API)
MathArena	matharena.ai	Live competition tracking	GitHub: eth-sri/matharena
BBH/BBEH	github.com/google-deepmi nd/bbeh	None (static public)	HF: BBEH dataset + GitHub

Table 11: Contamination-Resistant Reasoning Benchmarks Directory

8.3 Agentic and Tool-Calling Benchmarks

Benchmark	URL	MCP ?	Access Method
SWE-bench	swebench.com	No	HF: princeton-nlp/SWE-bench_Verified + GitHub
Terminal-Bench	tbench.ai	No	GitHub: harbor-framework/terminal-bench
MCP-Atlas	arxiv.org/abs/2602.00933	Yes	HF: papers/2602.00933
WebArena	webarena.dev	No	GitHub: web-arena-x/webarena
OSWorld	os-world.github.io	No	GitHub: xlang-ai/osworld + HF (migrating)
GAIA	arxiv.org/abs/2311.12983	No	HF: gaia-benchmark/GAIA (gated)
AgentBench	github.com/THUDM/AgentBench	No	GitHub: THUDM/AgentBench
BFCL v4	gorilla.cs.berkeley.edu	No	HF: gorilla-llm/Berkeley-Function-Calling-Leaderboard
ToolBench	github.com/openbmb/toolbench	No	HF: tuandunghcmut/toolbench-v1 + GitHub
tau-bench/tau2-bench	taubench.com	No	GitHub: sierra-research/tau-bench

MCP-Bench	github.com/Accenture/mcp-bench	Yes	GitHub: Accenture/mcp-bench
MCP-AgentBench	ojs.aaai.org (AAAI 2025)	Yes	Paper + evaluation framework
MCPVerse	arxiv.org/abs/2508.16260	Yes	arXiv:2508.16260
FeatureBench	github.com/LiberCoders/FeatureBench	No	GitHub: LiberCoders/FeatureBench
LongCLI-Bench	arxiv.org/abs/2602.14337	No	arXiv:2602.14337

Table 12: Agentic and Tool-Calling Benchmarks Directory

8.4 Efficiency and Cost Profiling Platforms

Platform	URL	Key Metrics	API Access
Artificial Analysis	artificialanalysis.ai	TPS, TTFT, pricing, intelligence	REST API (free)
HF LLM-Perf	huggingface.co/spaces/optimum/llm-perf-leaderboard	Throughput, latency, memory, energy	HF datasets
OpenRouter	openrouter.ai	Pricing, context, availability	REST API (free for models)
SemiAnalysis InferenceMAX	inference.semianalysis.com	TPS, TTFT, perf/\$, tokens/MW	Dashboard only
vLLM Benchmarks	docs.vllm.ai	TPS, TTFT, TPOT, E2E latency	CLI/Python (open source)
NVIDIA GenAI-Perf	docs.nvidia.com	TPS, TTFT, TPOT, P50/P90/P99	CLI tool (open source)
MLPerf Inference	mlcommons.org	Throughput, latency, multi-stream	Static results tables
Epoch AI	epoch.ai	Cost-per-performance-level, price trends	Web only
Modal LLM Almanac	modal.com/llm-almanac	TPS, TTFT, cost per 1M tokens	Interactive dashboard
BentoML llm-optimizer	github.com/bentoml/llm-optimizer	TPS, TTFT, P99, GPU utilization	CLI/Python (open source)

Table 13: Efficiency and Cost Profiling Platforms Directory

8.5 Recommended Data Aggregation Architecture

Based on the complete research findings, the optimal strategy for building a unified LLM evaluation data layer follows a four-tier architecture. Tier 1 provides real-time data via direct API integration with OpenRouter `/api/v1/models` (358+ models with pricing, context, and capabilities, no authentication

required for the models list), Artificial Analysis /api/* (speed, latency, intelligence scores, API key required, 1,000 requests per day free tier), and ZeroEval/LLM Stats api.zerogpt.com/v1/ (models with providers, features, and pricing, API key required). Tier 2 provides batch-refresh data via Hugging Face Hub datasets for the Open LLM Leaderboard results, SWE-bench Verified, GPQA (gated), GAIA (gated), BFCL, LMSYS Arena conversations, and ToolBench. Tier 3 provides periodic synchronization via GitHub scraping for ARC-AGI Community Leaderboard, AgentBench, WebArena, OSWorld, Terminal-Bench, and HELM. Tier 4 covers platforms with no programmatic access (Scale AI SEAL, LiveBench website) that require manual web scraping or direct partnership agreements.

For the highest-value automated monitoring pipeline, we recommend combining three real-time API sources: OpenRouter for immediate model availability and pricing, Artificial Analysis for quality and speed benchmarks, and ZeroEval for provider feature matrices. These three APIs together provide a comprehensive, machine-readable view of the LLM landscape that can be polled at hourly intervals to detect new model additions, pricing changes, and benchmark score updates.